

# Lab 2 - Spark & Big Data

...

Luca Telloli  
luca.telloli@upf.edu

Pablo Martín Castagnaro  
pablo.castagnaro@upf.edu

# Summary

- Considerations on Lab 1
- Spark
- Elastic Map Reduce (EMR)
- Running Spark on EMR
- Example: Running Word Count on EMR

# Learnings of Lab 1

- **Java**: review I/O (read & write files), basic Exceptions
- Parsing **JSON** to an instance of a Java Class and vice-versa
- **S3**:
  - with the **CLI** or **Web UI**: create buckets, change permissions
  - programmatically with the **SDK**: S3Uploader
-

# Considerations on Lab 1

- Read all data in a **single** node <- does not scale well
- Read all data **sequentially** <- multi-thread programming? difficult & still limited
- **Processing** is a small part of the program, still it's the **relevant** part
- I/O is a considerable part of the application and yet is not the focus, while being mostly tedious & repetitive

# Spark review

# What is an RDD?

- **Resilient Distributed Dataset:** Like a big table, broken into parts, each part available on a separate node
- It's an immutable distributed collection of objects
- Spark provides **two types** for such datasets

Resilient	Can recover from abnormal events (i.e.: hardware, network issues)
Distributed	Data might be held at different times in different storage elements (memory, disks, nodes), but it's still managed by the user as a single continuous sequence of items.

# What is an RDD?

RDDs can be of two families

They can contain any Java class (RDD<Int>, RDD<User>, etc...)

JavaRDD<V>	JavaPairRdd<K, V>
A dataset with a single type of value	A dataset with two elements, a key, and a value

Row 1

V1

Row 2

V2

Row n

Vn

K1

V1

K2

V2

Kn

Vn

# What is an RDD?

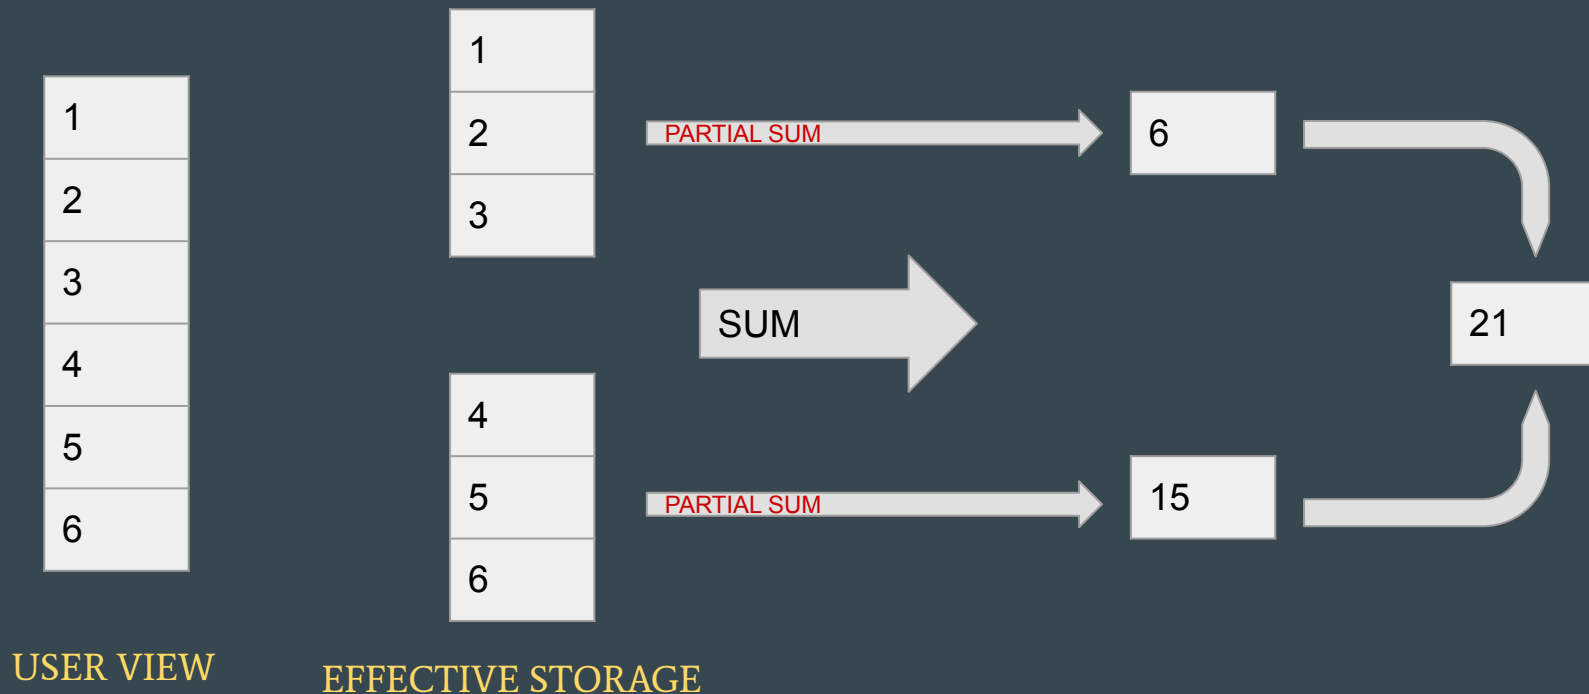
- Transformations on RDDs are defined differently depending on the type

	JavaRDD<V>	JavaPairRdd<K, V>
Mapping:	map()	mapValues()
Reducing:	reduce()	reduceByKey()
Flattening:	flatMap()	flatMapValues()



# What is an RDD?

Spark takes care of **orchestrating** the sequence of distributed operations



# Common spark functions

```
.map(elem-> ...)
```

Applies a transformation 1 to 1 to each element

```
.mapToPair(elem-> transformation(elem))
```

Applies a transformation 1 to 1 to each element. Generates an entry of a PairRDD for each of them

```
.flatMap(elem-> transformation(elem))
```

Applies a transformation 1 to N to each element. **If the resulting type is a list, it requires an iterator** (see WordCount example for that)

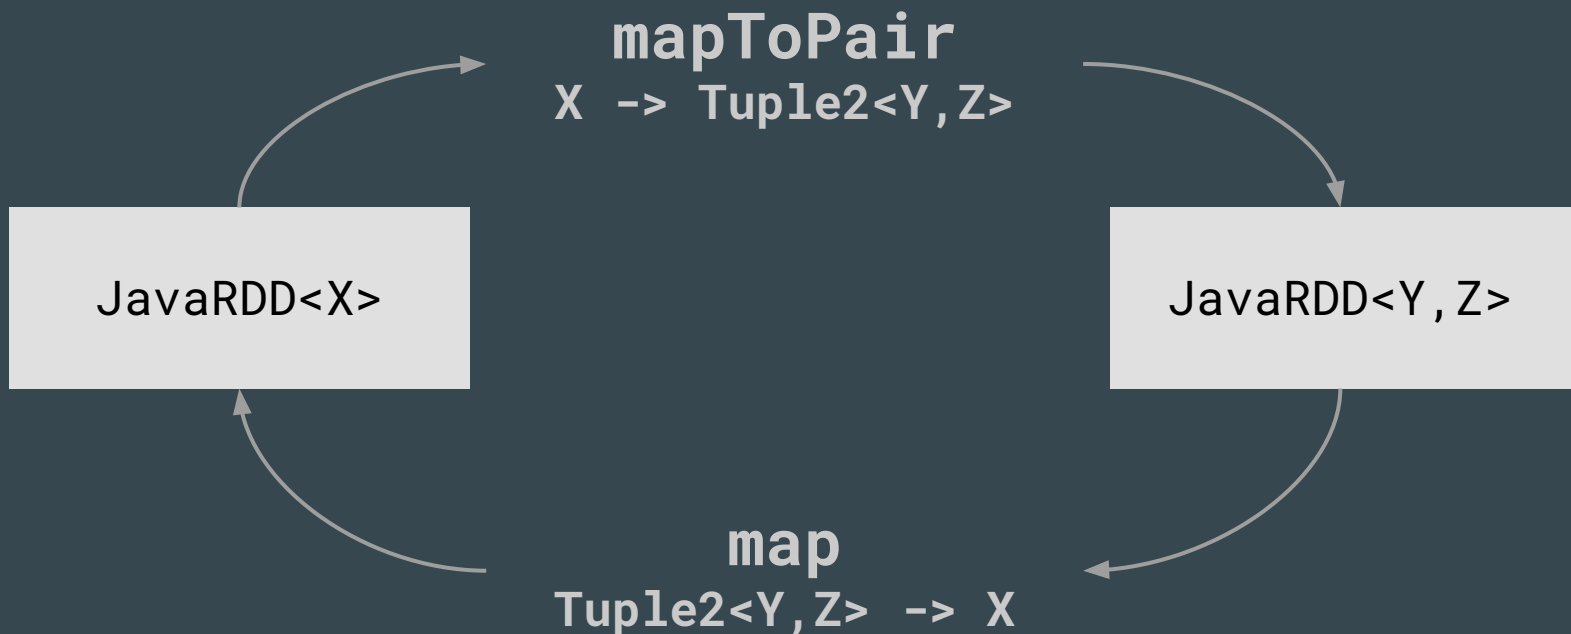
```
.filter(elem->condition(elem))
```

Selects all the elements satisfying a condition.

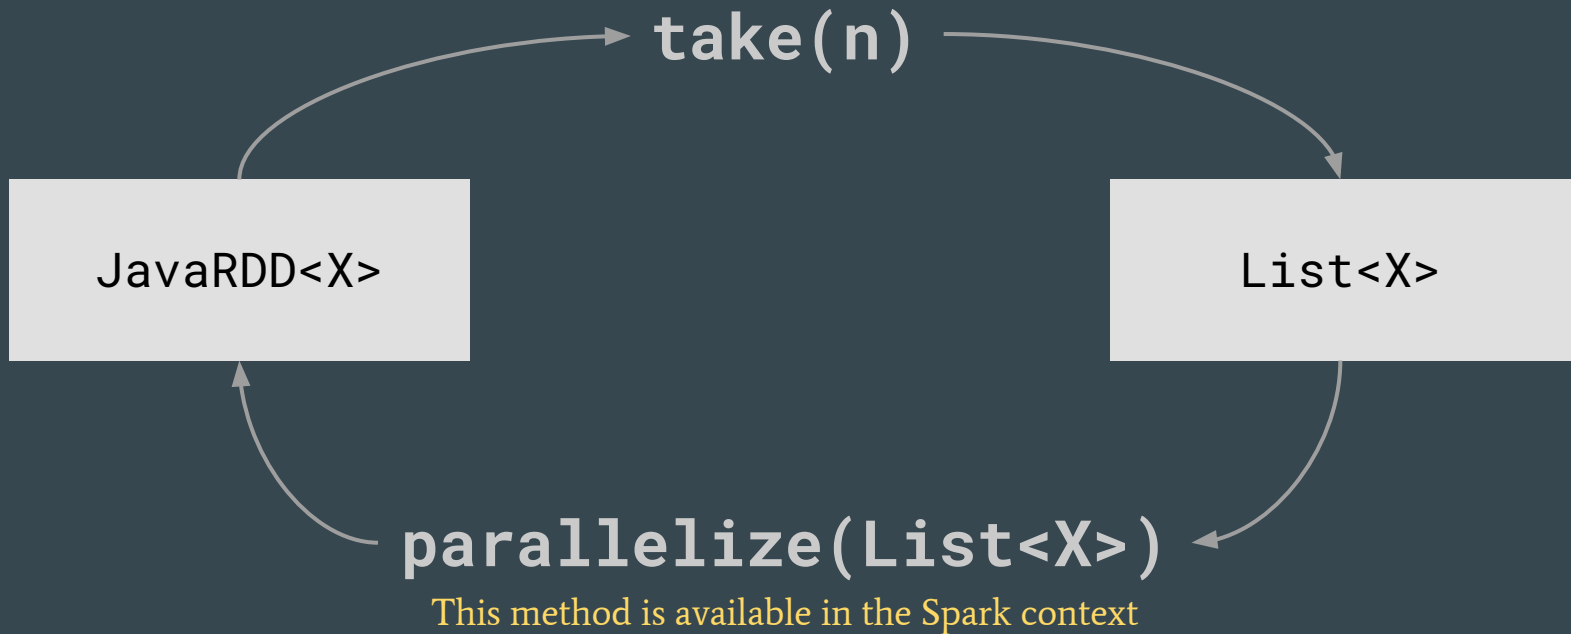
```
.reduceByKey(key1, key2->...)
```

Groups and transforms all the values from rows having identical keys (requires a PairRDD)

# Moving between $\text{RDD}\langle K \rangle$ and $\text{RDD}\langle K, V \rangle$



# Moving between RDDs and Lists



# Common spark functions

JavaRDD<V>	JavaPairRDD<K,V>
<code>sort([ascending=True])</code>	<code>sortByKey([ascending=True])</code>
<code>take()</code> Materializes N elements from a given RDD	<code>take(N)</code> Materializes N elements from a given PairRDD
<code>join()</code>	<code>join()</code> Returns a dataset of (K, (V, W)) pairs with all pairs of elements for each key.

Items in a JavaPairRDD are of type [scala.Tuple2](#)

Spark does not have a `sortByValue()` function. If you want to do so, you need to swap ((K,V) to (V,K)) tuples in the RDD, `sortByKey` and swap again. See the function `swap()` of the `Tuple2` class.

# Review

- Functional operators in Spark: continuous transformation of a  $RDD<T>$  to a  $RDD<U>$  using functional operators
  - Transformations & Actions
- Spark execution model: only Actions operation trigger a materialisation (data physically present on a storage level).
- Programming guide:  
<https://spark.apache.org/docs/2.4.7/rdd-programming-guide.html>

# Installing Spark locally

If you install Spark locally (*highly recommended*), please **use the latest stable version**.

Check your installed version with **spark-submit --version**

# Spark Applications in Java

- A Spark application is a Java application that uses the Apache Spark libraries
- In a new Spark application, create first a `SparkConfiguration` and pass it to a `JavaSparkContext`, as displayed in this example:

```
SparkConf conf = new SparkConf().setAppName(<your app name>);  
  
JavaSparkContext sc = new JavaSparkContext(conf);
```

An example application for **WordCount** is available on Moodle, and can be used a starting point for your applications.



# Running spark applications

It's frequent during the development phase to run a Spark application locally, usually on a smaller subset of the input.

To run the application locally, you **won't use Java**; instead, you'll use the Spark framework, accessible with the **spark-submit** command. The command will be similar to:

```
spark-submit [--master <MASTER>] --class <MAIN CLASS> your.jar arg1 ... argN
```

# Example

# Running WordCount locally

Download it from moodle

# Elastic Map Reduce (EMR)


# What is EMR ?

- AWS EMR is a tool for big data processing and analysis
- Supports processing of large datasets in a **distributed computing environment**
- AWS EMR offers expandable, low-configuration, managed service
- Based on Apache Hadoop (MapReduce) cluster of virtual servers on AWS EC2 and AWS S3

# Running Spark on EMR

# Access to **AWS Academy** and its resources

0. Check your email and register in AWS Academy with the provided link



1. Click on **Courses** to access the AWS resources


2. Click on **Foundational Services** and access terms and conditions

# Access to **AWS Academy** and its resources

3. Click on **Start Lab** to access the AWS resources

4. Wait until **your lab** is ready

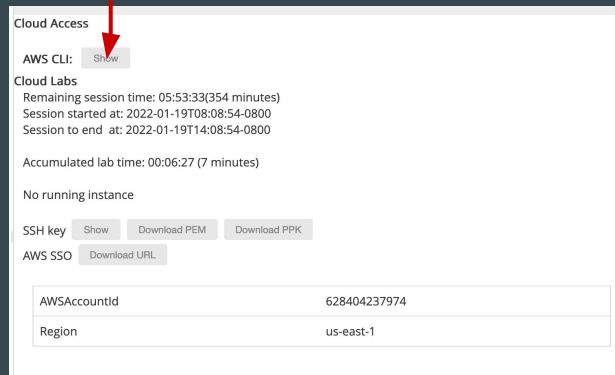
5. Click **AWS Details** to access your account credentials



# What does AWS Academy offer?

- “Student” version of the AWS console
- Pre-loaded with some expendable credit (\$100)
- The CLI credentials expire every 3 hours

6. Click on **Show** to get CLI Credentials



The screenshot shows the 'Cloud Access' section of the AWS Academy interface. It includes a 'Show' button for the 'AWS CLI' section. Below this, the 'Cloud Labs' section displays session information: 'Remaining session time: 05:53:33(354 minutes)', 'Session started at: 2022-01-19T08:08:54-0800', and 'Session to end at: 2022-01-19T14:08:54-0800'. It also shows 'Accumulated lab time: 00:06:27 (7 minutes)' and 'No running instance'. At the bottom, there are buttons for 'Show', 'Download PEM', and 'Download PPK' for the 'SSH key', and a 'Download URL' button for 'AWS SSO'. A table at the bottom lists the 'AWSAccountId' as '628404237974' and the 'Region' as 'us-east-1'.

AWSAccountId	628404237974
Region	us-east-1

7. Click on **AWS Console** to open the Web Console






# AWS console

7. Click on **AWS Console**  
to open the Web Console



# Setting up a Cluster on AWS

- Login into the AWS console through your Educate account
- Search or browse for the EMR service (Elastic Map Reduce)
- Select: *Create Cluster*
- It's better to create your first cluster immediately, because the first run can take longer





Services ▾

Resource Groups ▾



vocstartsof

## Create Cluster - Quick Options [Go to advanced options](#)

### General Configuration

Cluster name

☒ Logging ⓘ

S3 folder

Launch mode ☒ Cluster ⓘ ☐ Step execution ⓘ

### Software configuration

Release

Applications ☒ Core Hadoop: Hadoop 2.8.4 with Ganglia 3.7.2, Hive 2.3.3, Hue 4.2.0, Mahout 0.13.0, Pig 0.17.0, and Tez 0.8.4

☐ HBase: HBase 1.4.6 with Ganglia 3.7.2, Hadoop 2.8.4, Hive 2.3.3, Hue 4.2.0, Phoenix 4.14.0, and ZooKeeper 3.4.12

☐ Presto: Presto 0.206 with Hadoop 2.8.4 HDFS and Hive 2.3.3 Metastore

☐ Spark: Spark 2.3.1 on Hadoop 2.8.4 YARN with Ganglia 3.7.2 and Zeppelin 0.7.3

☐ Use AWS Glue Data Catalog for table metadata ⓘ

1. Give it a name

2. Use emr-5.17.2

3. Select Spark support

### Hardware configuration

Instance type

m4.large

The selected instance type adds a default 32 GiB GP2 EBS volume per instance. [Learn more](#)

Number of instances

3

(1 master and 2 core nodes)

### Security and access

EC2 key pair

No key pairs found

[Learn how to create an EC2 key pair.](#)

Permissions

☒ Default ☐ Custom

Use default IAM roles. If roles are not present, they will be automatically created for you with managed policies for automatic policy updates.

EMR role

[EMR\\_DefaultRole](#)



EC2 instance profile

[EMR\\_EC2\\_DefaultRole](#)




Cancel

Create cluster

4. Select m4.large instances

5. If you want, you can configure here an EC2 key pair to access other interfaces with more information about your cluster

6. Press *Create Cluster*



You'll find yourself here

Your cluster is currently provisioning, it could take some minutes and maybe the first time it might need to wait for an authorisation: it should be a matter of minutes

The screenshot shows the AWS EMR console interface. At the top, there's a navigation bar with the AWS logo, 'Services', 'Resource Groups', and a user profile. The left sidebar contains a menu with 'Amazon EMR' and various sub-features like 'Clusters', 'Security configurations', 'VPC subnets', 'Events', 'Notebooks', 'Help', and 'What's new'. The main content area displays details for a cluster named 'My cluster', which is in a 'Waiting' state. Above the cluster name are buttons for 'Clone', 'Terminate', and 'AWS CLI export'. Below the cluster name are tabs for 'Summary', 'Application history', 'Monitoring', 'Hardware', 'Events', 'Steps', 'Configurations', and 'Bootstrap actions'. The 'Steps' tab is currently selected. In the 'Steps' section, there's a blue 'Add step' button and two grey buttons: 'Cancel step' and 'Delete step'. Below this is a table of steps. The table has columns for 'Filter', 'ID', 'Name', 'Status', 'Start time (UTC+1)', 'Elapsed time', and 'Log files'. One step is listed with ID 's-314T9KZT8BP3T', Name 'Setup hadoop debugging', and Status 'Pending'. A red arrow points from the 'Add step' button to the 'Steps' tab, and another red arrow points from the 'Steps' tab to the pending step in the table.

Cluster: My cluster **Waiting** Cluster ready after last step completed.

Summary Application history Monitoring Hardware Events **Steps** Configurations Bootstrap actions

**Add step** Cancel step Delete step

Filter: All steps Filter steps ... 1 step (all loaded) View all interactive jobs View all jobs

ID	Name	Status	Start time (UTC+1)	Elapsed time	Log files
s-314T9KZT8BP3T	Setup hadoop debugging	Pending		--	View logs

After bootstrapping the cluster is green and waiting to run applications.

**Remember to stop it after your steps end!**

Select the *Steps* tab

Select *Add step* to configure a Spark application

**Add step**

**Step type**  
Custom JAR  
Streaming program  
✓ Spark application

**Name**  
Spark application

**Deploy mode**  
Cluster  
Run your driver on a slave node (cluster mode) or on the master node as an external client (client mode).  
Specify other options for spark-submit.

**Spark-submit options**

**Application location\***  
s3://  
Path to a JAR with your application and dependencies (client deploy mode only supports a local path).  
Specify optional arguments for your application.

**Arguments**

**Action on failure**  
Continue  
What to do if the step fails.

Cancel Add

Select *Spark Application*

Provide a name

The options are the same that you'd give to `spark-submit` on the command line

Location on S3 where you put your JAR

The arguments are the same that you'd give on the command line

Go!

## Example of parameters configuration to execute the WordCount

Añadir paso

X

Tipo de paso

Aplicación de Spark

▼

Nombre

Aplicación de Spark

Modo implementación

Clúster

▼

Ejecute el controlador en un nodo esclavo (modo clúster) o en el nodo maestro como cliente externo (modo cliente).

Opciones de spark-submit

--class spark.WordCount

Especifique otras opciones para spark-submit.

Ubicación de la aplicación\*

s3://test.lsd.2022.david/jars/spark-test-1.0-SNAF

📁

La ruta al archivo JAR con su aplicación y dependencias (el modo implementación del cliente solo es compatible con la ruta local).

Argumentos

s3://test.lsd.2022.david/input  
/shakespeare.txt  
s3://test.lsd.2022.david/output  
/execution1

Especifique los argumentos opcionales para su aplicación.

Acción sobre el error

Continuar

▼

Qué hacer si se produce un error en el paso.

Cancelar

Añadir



```
luke@rameau:$ aws s3 ls  
s3://lsds2018-lab2/output/
```

2019-02-04	20:56:03	0	_SUCCESS
2019-02-04	20:56:02	25436	part-00000
2019-02-04	20:56:02	25185	part-00001
2019-02-04	20:56:02	25724	part-00002
2019-02-04	20:56:02	25239	part-00003
2019-02-04	20:56:03	25095	part-00004
2019-02-04	20:56:03	25402	part-00005
2019-02-04	20:56:03	25684	part-00006
2019-02-04	20:56:03	25083	part-00007
2019-02-04	20:56:03	25896	part-00008
2019-02-04	20:56:03	26084	part-00009
2019-02-04	20:56:03	24749	part-00010
2019-02-04	20:56:03	25763	part-00011
2019-02-04	20:56:03	25568	part-00012
2019-02-04	20:56:03	25193	part-00013
2019-02-04	20:56:03	25865	part-00014
2019-02-04	20:56:03	26138	part-00015

This is a typical output from an hadoop job, notice the `_SUCCESS` empty file that denotes the success of the execution

Example

Running WordCount on EMR

## Part 1: Run the application locally



Download the project “WordCount Spark Example” from AulaGlobal.


Use **mvn package** to generate an executable JAR from it

Download Shakespeare’s work from [here](#)

Use **spark-submit** to run locally the generated JAR for the downloaded data

How many times does the word ‘trojan’ appear?

## Part 2: Now on EMR



- Create a bucket for the experiment
- Upload the generated jar and the text file
- Go to your EMR cluster:
  - 1) Add step
    - a) Set JAR path to S3
    - b) Set arguments
      - i) Input file path
      - ii) Output file path
  - 2) Wait for execution

Can you see your results on S3?